

```

import subprocess

process = subprocess.Popen(["python", "-c", "import imbalanced_learn;
print(imbalanced_learn.__file__)"], stdout=subprocess.PIPE)
output, err = process.communicate()

if err:
    print("Error:", err)
else:
    print("imbalanced_learn.__file__:", output.decode("utf-8").strip())

imbalanced_learn.__file__:
rm -rf <path_to_imblearn_dir>

/bin/bash: -c: line 1: syntax error near unexpected token `newline'
/bin/bash: -c: line 1: `rm -rf <path_to_imblearn_dir>'

import subprocess

process = subprocess.Popen(["python", "-c", "import sklearn;
print(sklearn.__file__)"], stdout=subprocess.PIPE)
output, err = process.communicate()

if err:
    print("Error:", err)
else:
    print("sklearn.__file__:", output.decode("utf-8").strip())

sklearn.__file__:
/usr/local/lib/python3.10/dist-packages/sklearn/__init__.py

rm -rf <path_to_sklearn_dir>

/bin/bash: -c: line 1: syntax error near unexpected token `newline'
/bin/bash: -c: line 1: `rm -rf <path_to_sklearn_dir>'

!pip cache purge

Files removed: 4

!pip install scikit-learn
!pip install imblearn

Requirement already satisfied: scikit-learn in
/usr/local/lib/python3.10/dist-packages (1.3.2)
Requirement already satisfied: numpy<2.0,>=1.17.3 in
/usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.23.5)
Requirement already satisfied: scipy>=1.5.0 in
/usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.11.3)
Requirement already satisfied: joblib>=1.1.1 in

```

```

/usr/local/lib/python3.10/dist-packages (from scikit-learn) (1.3.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/usr/local/lib/python3.10/dist-packages (from scikit-learn) (3.2.0)
Requirement already satisfied: imblearn in
/usr/local/lib/python3.10/dist-packages (0.0)
Requirement already satisfied: imbalanced-learn in
/usr/local/lib/python3.10/dist-packages (from imblearn) (0.10.1)
Requirement already satisfied: numpy>=1.17.3 in
/usr/local/lib/python3.10/dist-packages (from imbalanced-learn-
>imblearn) (1.23.5)
Requirement already satisfied: scipy>=1.3.2 in
/usr/local/lib/python3.10/dist-packages (from imbalanced-learn-
>imblearn) (1.11.3)
Requirement already satisfied: scikit-learn>=1.0.2 in
/usr/local/lib/python3.10/dist-packages (from imbalanced-learn-
>imblearn) (1.3.2)
Requirement already satisfied: joblib>=1.1.1 in
/usr/local/lib/python3.10/dist-packages (from imbalanced-learn-
>imblearn) (1.3.2)
Requirement already satisfied: threadpoolctl>=2.0.0 in
/usr/local/lib/python3.10/dist-packages (from imbalanced-learn-
>imblearn) (3.2.0)

```

```
cd /content/drive/MyDrive/StrokeProject
```

```
/content/drive/MyDrive/StrokeProject
```

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import missingno as ms
import sklearn

```

```

train = pd.read_csv('train_2v.csv')
test = pd.read_csv('test_2v.csv')
train.head()

```

| | id | gender | age | hypertension | heart_disease | ever_married | \ |
|---|-------|--------|------|--------------|---------------|--------------|---|
| 0 | 30669 | Male | 3.0 | 0 | 0 | No | |
| 1 | 30468 | Male | 58.0 | 1 | 0 | Yes | |
| 2 | 16523 | Female | 8.0 | 0 | 0 | No | |
| 3 | 56543 | Female | 70.0 | 0 | 0 | Yes | |
| 4 | 46136 | Male | 14.0 | 0 | 0 | No | |

| | work_type | Residence_type | avg_glucose_level | bmi | smoking_status | \ |
|---|-----------|----------------|-------------------|------|----------------|---|
| 0 | children | Rural | 95.12 | 18.0 | NaN | |
| 1 | Private | Urban | 87.96 | 39.2 | never | |

```

smoked
2      Private      Urban      110.89  17.6
NaN
3      Private      Rural      69.04  35.9  formerly
smoked
4  Never_worked      Rural      161.28  19.1
NaN

      stroke
0         0
1         0
2         0
3         0
4         0

```

In the dataset we have 12 columns where 11 contains the features and the last one contains the result

```

test.head()

   id  gender  age  hypertension  heart_disease  ever_married  \
0  36306   Male  80.0           0             0             Yes
1  61829  Female  74.0           0             1             Yes
2  14152  Female  14.0           0             0             No
3  12997   Male  28.0           0             0             No
4  40801  Female  63.0           0             0             Yes

      work_type  Residence_type  avg_glucose_level  bmi
smoking_status
0      Private      Urban      83.84  21.1  formerly
smoked
1  Self-employed      Rural      179.50  26.0  formerly
smoked
2      children      Rural      95.16  21.2
NaN
3      Private      Urban      94.76  23.4
NaN
4  Govt_job      Rural      83.57  27.6  never
smoked

train.shape
(43400, 12)

test.shape
(18601, 11)

```

Data Cleaning Identifying missing *attributes*

```
train_missing_values=train.isnull().sum()  
train_missing_values
```


```
id                0  
gender            0  
age              0  
hypertension      0  
heart_disease     0  
ever_married      0  
work_type         0  
Residence_type    0  
avg_glucose_level 0  
bmi              1462  
smoking_status    13292  
stroke            0  
dtype: int64
```

```
test_missing_values=test.isnull().sum()  
test_missing_values
```

```
id                0  
gender            0  
age              0  
hypertension      0  
heart_disease     0  
ever_married      0  
work_type         0  
Residence_type    0  
avg_glucose_level 0  
bmi              591  
smoking_status    5751  
dtype: int64
```

```
ms.matrix(train)
```

```
<Axes: >
```




Removing missing value

```
train_data=train.dropna(axis=0,how="any")
test_data=test.dropna(axis=0,how="any")
print('train data shape: {}'.format(train_data.shape))
print('test data shape: {}'.format(test_data.shape))
```

```
train data shape: (29072, 12)
test data shape: (12423, 11)
```


```
ms.matrix(train_data)
```

```
<Axes: >
```



```
ms.matrix(test_data)
```

```
<Axes: >
```



Pattern Recognition

```
train_data["stroke"].value_counts()
```

```
0    28524
```


```
1     548
```

```
Name: stroke, dtype: int64
```

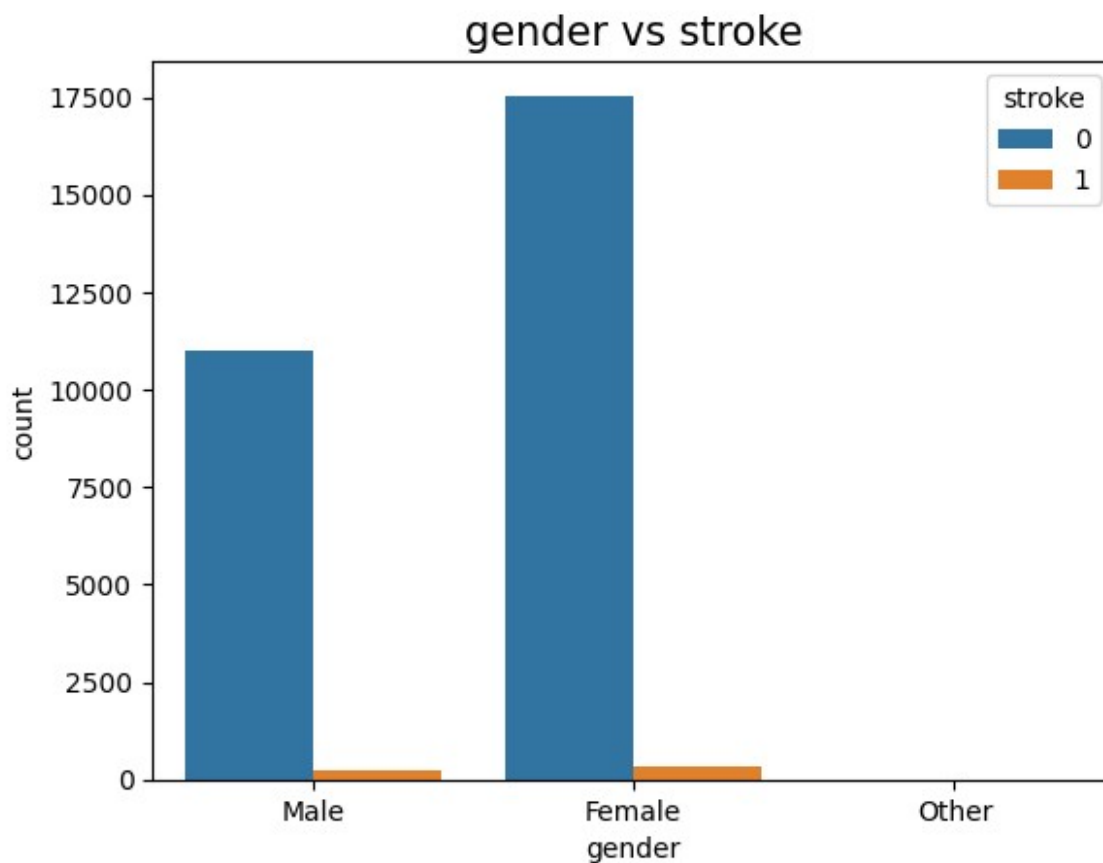
```
sns.countplot(x=train_data["stroke"])
```

```
plt.title("no of patients affected by stroke", fontsize=15)
```

```
plt.show()
```



```
sns.countplot(x=train_data["gender"], hue=train_data["stroke"])  
plt.title("gender vs stroke", fontsize=15)  
plt.show()
```



```
train_data.groupby(["gender"])[ "stroke" ].value_counts()
```

```
gender  stroke
Female  0      17539
        1        313
Male    0     10978
        1        235
Other   0         7
Name: stroke, dtype: int64
```

```
train_data["smoking_status"].value_counts()
```

```
never smoked    15747
formerly smoked    7099
smokes          6226
Name: smoking_status, dtype: int64
```


```
train_data.groupby(["smoking_status"])[ "stroke" ].value_counts()
```

```
smoking_status  stroke
formerly smoked  0      6919
                1       180
never smoked    0     15491
                1       256
```



```
smokes      0      6114
           1      112
Name: stroke, dtype: int64

sns.countplot(x=train_data["gender"],
hue=train_data["smoking_status"])
plt.title("gender vs type of smokers", fontsize=15)
plt.show()
```



Encoding data

```
str_data=train_data.select_dtypes(include=['object'])
str_dt=test_data.select_dtypes(include=['object'])

int_data=train_data.select_dtypes(include=['integer', 'float'])
int_dt=test_data.select_dtypes(include=['integer', 'float'])

from sklearn.preprocessing import LabelEncoder

label=LabelEncoder()
features=str_data.apply(label.fit_transform)
features=features.join(int_data)
features.head()
```

| id \ | gender | ever_married | work_type | Residence_type | smoking_status |
|---------|--------|--------------|-----------|----------------|----------------|
| 1 30468 | 1 | 1 | 2 | 1 | 1 |
| 3 56543 | 0 | 1 | 2 | 0 | 0 |
| 6 52800 | 0 | 1 | 2 | 1 | 0 |
| 7 41413 | 0 | 1 | 3 | 0 | 1 |
| 8 15266 | 0 | 1 | 2 | 0 | 2 |

| | age | hypertension | heart_disease | avg_glucose_level | bmi | stroke |
|---|------|--------------|---------------|-------------------|------|--------|
| 1 | 58.0 | 1 | 0 | 87.96 | 39.2 | 0 |
| 3 | 70.0 | 0 | 0 | 69.04 | 35.9 | 0 |
| 6 | 52.0 | 0 | 0 | 77.59 | 17.7 | 0 |
| 7 | 75.0 | 0 | 1 | 243.53 | 27.0 | 0 |
| 8 | 32.0 | 0 | 0 | 77.67 | 32.3 | 0 |

```
test1=str_dt.apply(label.fit_transform)
Test=test1.join(int_dt)
Test.head()
```

| id \ | gender | ever_married | work_type | Residence_type | smoking_status |
|---------|--------|--------------|-----------|----------------|----------------|
| 0 36306 | 1 | 1 | 2 | 1 | 0 |
| 1 61829 | 0 | 1 | 3 | 0 | 0 |
| 4 40801 | 0 | 1 | 0 | 0 | 1 |
| 5 9348 | 0 | 1 | 2 | 1 | 1 |
| 7 60512 | 1 | 1 | 0 | 1 | 1 |

| | age | hypertension | heart_disease | avg_glucose_level | bmi |
|---|------|--------------|---------------|-------------------|------|
| 0 | 80.0 | 0 | 0 | 83.84 | 21.1 |
| 1 | 74.0 | 0 | 1 | 179.50 | 26.0 |
| 4 | 63.0 | 0 | 0 | 83.57 | 27.6 |
| 5 | 66.0 | 1 | 0 | 219.98 | 32.2 |
| 7 | 46.0 | 0 | 0 | 120.80 | 32.5 |

Modeling & predicting the data

```
xtrain=features.drop(["stroke"],axis=1)
xtrain.shape
```

```
(29072, 11)
```

```
ytrain=features["stroke"]
ytrain.head()
ytrain.shape
```

```
(29072,)
```

```
from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test=train_test_split(xtrain, ytrain)
```

```
x_test.shape
```

```
(7268, 11)
```

```
y_test.shape
```

```
(7268,)
```

```
x_train.head()
```

| | gender | ever_married | work_type | Residence_type | smoking_status |
|-------|--------|--------------|-----------|----------------|----------------|
| id \ | | | | | |
| 14628 | 0 | 0 | 2 | 0 | 1 |
| 33280 | | | | | |
| 17895 | 1 | 1 | 3 | 1 | 2 |
| 70939 | | | | | |
| 10614 | 1 | 1 | 2 | 0 | 0 |
| 48669 | | | | | |
| 34892 | 1 | 1 | 2 | 1 | 2 |
| 51685 | | | | | |
| 28149 | 0 | 1 | 3 | 0 | 0 |
| 40489 | | | | | |

| | age | hypertension | heart_disease | avg_glucose_level | bmi |
|-------|------|--------------|---------------|-------------------|------|
| 14628 | 37.0 | 0 | 0 | 93.80 | 24.6 |
| 17895 | 61.0 | 0 | 0 | 75.28 | 19.6 |
| 10614 | 71.0 | 0 | 0 | 66.58 | 27.5 |
| 34892 | 44.0 | 0 | 0 | 228.40 | 36.1 |
| 28149 | 59.0 | 0 | 0 | 82.92 | 34.9 |

```
y_train.head()
```

| | |
|-------|---|
| 14628 | 0 |
| 17895 | 0 |
| 10614 | 0 |
| 34892 | 0 |

```
28149    0
Name: stroke, dtype: int64
```

Naive Bayes

```
x_test.head()
```

| | gender | ever_married | work_type | Residence_type | smoking_status |
|-------|--------|--------------|-----------|----------------|----------------|
| id \ | | | | | |
| 805 | 1 | 0 | 2 | 0 | 1 |
| 58037 | | | | | |
| 31952 | 1 | 0 | 2 | 1 | 0 |
| 62918 | | | | | |
| 5083 | 1 | 1 | 3 | 0 | 0 |
| 3820 | | | | | |
| 12863 | 0 | 1 | 2 | 0 | 0 |
| 14684 | | | | | |
| 21818 | 1 | 1 | 2 | 0 | 1 |
| 63056 | | | | | |

| | age | hypertension | heart_disease | avg_glucose_level | bmi |
|-------|------|--------------|---------------|-------------------|------|
| 805 | 21.0 | 0 | 0 | 78.52 | 27.2 |
| 31952 | 40.0 | 0 | 0 | 63.29 | 29.7 |
| 5083 | 77.0 | 0 | 0 | 88.75 | 31.9 |
| 12863 | 65.0 | 0 | 0 | 244.95 | 36.1 |
| 21818 | 78.0 | 1 | 0 | 90.04 | 25.6 |

```
y_test.head()
```

```
805    0
31952   0
5083    0
12863   0
21818   0
Name: stroke, dtype: int64
```

```
from sklearn.naive_bayes import GaussianNB
```

```
model=GaussianNB()
model.fit(x_train, y_train)
```

```
GaussianNB()
```

```
predict=model.predict(x_test)
predict
```

```
array([0, 0, 0, ..., 0, 0, 0])
```

```
test_score=model.score(x_test, y_test)
print("NBtest_score:", test_score)
```

NBtest_score: 0.9744083654375344

```
nb_conf_mtr=pd.crosstab(y_test, predict)
nb_conf_mtr
```

```
col_0      0      1
stroke
0         7077    61
1          125     5
```

```
from sklearn.metrics import classification_report
```

```
nbreport=classification_report(y_test, predict)
print(nbreport)
```

| | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0 | 0.98 | 0.99 | 0.99 | 7138 |
| 1 | 0.08 | 0.04 | 0.05 | 130 |
| accuracy | | | 0.97 | 7268 |
| macro avg | 0.53 | 0.51 | 0.52 | 7268 |
| weighted avg | 0.97 | 0.97 | 0.97 | 7268 |

Decision Tree

```
from sklearn.tree import DecisionTreeClassifier
```

```
dt_mod=DecisionTreeClassifier(max_depth=8)
dt_mod.fit(x_train, y_train)
```

```
DecisionTreeClassifier(max_depth=8)
```

```
y_predict=dt_mod.predict(x_test)
y_predict
```

```
array([0, 0, 0, ..., 0, 0, 0])
```

```
ts_dt_score=dt_mod.score(x_test, y_test)
print("Decision tree test score:", ts_dt_score)
```

Decision tree test score: 0.9794991744634012

```
dectree_report=classification_report(y_test, y_predict)
print(dectree_report)
```

| | precision | recall | f1-score | support |
|---|-----------|--------|----------|---------|
| 0 | 0.98 | 1.00 | 0.99 | 7138 |
| 1 | 0.12 | 0.02 | 0.04 | 130 |

| | | | | |
|--------------|------|------|------|------|
| accuracy | | | 0.98 | 7268 |
| macro avg | 0.55 | 0.51 | 0.51 | 7268 |
| weighted avg | 0.97 | 0.98 | 0.97 | 7268 |

```
dt_conf_mtr=pd.crosstab(y_test, y_predict)
dt_conf_mtr
```

| | | |
|--------|------|----|
| col_0 | 0 | 1 |
| stroke | | |
| 0 | 7116 | 22 |
| 1 | 127 | 3 |

Random Forest

```
from sklearn.ensemble import RandomForestClassifier
rfc = RandomForestClassifier(n_estimators= 100)

rfc.fit(x_train,y_train)

RandomForestClassifier()

y_pred_rfc = rfc.predict(x_test)

print(classification_report(y_test,y_pred_rfc))
```

| | | | | |
|---|-----------|--------|----------|---------|
| | precision | recall | f1-score | support |
| 0 | 0.98 | 1.00 | 0.99 | 7138 |
| 1 | 0.00 | 0.00 | 0.00 | 130 |

| | | | | |
|--------------|------|------|------|------|
| accuracy | | | 0.98 | 7268 |
| macro avg | 0.49 | 0.50 | 0.50 | 7268 |
| weighted avg | 0.96 | 0.98 | 0.97 | 7268 |

```
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1471: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1471: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1471: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
```

```
frst_cnf= pd.crosstab(y_test,y_pred_rfc)
print(frst_cnf)
```

```
col_0    0
stroke
0       7138
1       130
```

Multi-Layer Perceptron Classifier

```
from sklearn.neural_network import MLPClassifier
mlp=MLPClassifier()
mlp.fit(x_train,y_train)
y_pred_mlp = mlp.predict(x_test)
mlp.score(x_test,y_test)
0.9821133736929004

from sklearn.model_selection import cross_val_score
cross_val_score(model,xtrain,ytrain,cv = 20,
scoring='accuracy').mean()
0.9757842475511938

cross_val_score(dt_mod,xtrain,ytrain,cv = 20,
scoring='accuracy').mean()
0.9790521152934069

cross_val_score(rfc,xtrain,ytrain,cv = 20, scoring='accuracy').mean()
0.9811503212534707

cross_val_score(mlp,xtrain,ytrain,cv = 20, scoring='accuracy').mean()
0.9392994241388353
```

Applying PCA

```
from sklearn.decomposition import PCA
pca = PCA(n_components=3)
principalComponents = pca.fit_transform(xtrain)

from sklearn.model_selection import train_test_split
x_train, x_test, y_train, y_test=train_test_split(xtrain, ytrain)

model_2=GaussianNB()
model_2.fit(x_train, y_train)
```

```

GaussianNB()

test_score=model_2.score(x_test, y_test)
print("NBtest_score:", test_score)

NBtest_score: 0.9753714914694551

dt_mod=DecisionTreeClassifier()
dt_mod.fit(x_train, y_train)

DecisionTreeClassifier()

ts_dt_score=dt_mod.score(x_test, y_test)
print("Decision tree test score:", ts_dt_score)

Decision tree test score: 0.9595487066593286

rfc.fit(x_train,y_train)

RandomForestClassifier()

y_pred_rfc = rfc.predict(x_test)

print(pd.crosstab(y_test,y_pred_rfc))
print(classification_report(y_test,y_pred_rfc))

```

| | | | | | |
|--------|--------------|-----------|--------|----------|---------|
| col_0 | 0 | | | | |
| stroke | | | | | |
| 0 | 7132 | | | | |
| 1 | 136 | | | | |
| | | precision | recall | f1-score | support |
| | 0 | 0.98 | 1.00 | 0.99 | 7132 |
| | 1 | 0.00 | 0.00 | 0.00 | 136 |
| | accuracy | | | 0.98 | 7268 |
| | macro avg | 0.49 | 0.50 | 0.50 | 7268 |
| | weighted avg | 0.96 | 0.98 | 0.97 | 7268 |

```

/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1471: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1471: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
/usr/local/lib/python3.10/dist-packages/sklearn/metrics/_classification.py:1471: UndefinedMetricWarning: Precision and F-score are ill-

```



```
defined and being set to 0.0 in labels with no predicted samples. Use  
`zero_division` parameter to control this behavior.  
_warn_prf(average, modifier, msg_start, len(result))
```

```
rfc.score(x_test,y_test)
```

```
0.9812878370941112
```

```
mlp=MLPClassifier()
```

```
mlp.fit(x_train,y_train)
```

```
y_pred_mlp = mlp.predict(x_test)
```

```
mlp.score(x_test,y_test)
```

```
0.9812878370941112
```

```
cross_val_score(model_2,xtrain,ytrain,cv = 20,  
scoring='accuracy').mean()
```

```
0.9757842475511938
```

```
cross_val_score(dt_mod,xtrain,ytrain,cv = 20,  
scoring='accuracy').mean()
```

```
0.9581380978121441
```

```
cross_val_score(rfc,xtrain,ytrain,cv = 20, scoring='accuracy').mean()
```

```
0.9811503212534707
```

```
cross_val_score(mlp,xtrain,ytrain,cv = 20, scoring='accuracy').mean()
```

```
0.9810127696716275
```